

MAKE 10 / LABOR & AI

THE HUMANS BEHIND THE MACHINE

WHO PAYS THE HUMAN COST OF MAKING AI INTELLIGENT?

A visual investigation of Kenyan data annotators, outsourced work, and the human judgments behind AI safety.

A CASE STUDY OF THE 2021-2022 OPENAI-SAMA DATA-LABELING PROJECT

AI DOES NOT WORK ALONE.

The polished chatbot interface hides a chain of decisions, contracts, and human review.

01 OPENAI / AI DEVELOPER

Commissioned outsourced data-labeling work to support a toxic-content detection tool.

↓

02 SAMA / SERVICE PROVIDER + EMPLOYER

Organized the project and employed the Kenyan workers who performed the labeling.

↓

03 KENYAN ANNOTATORS / HUMAN REVIEW

Read text passages and classified harmful material using human judgment.

↓

04 AI SAFETY TOOL / TECHNICAL OUTPUT

The labels helped develop a toxic-content detection tool that TIME reported was later integrated into ChatGPT.

This is a simplified view of the documented outsourcing relationship, not a complete technical diagram of ChatGPT training.

THE INVISIBLE STEP IS HUMAN WORK.

When the system works, users see an answer. They rarely see the people whose judgments helped make the system safer.

BEHIND EVERY LABEL IS A HUMAN DECISION.

The Kenyan team reviewed text excerpts and sorted harmful material into categories for a safety project.

~36 **WORKERS**
Approximate team size reported by TIME

3 **TEAMS**
Organized for the project

VIOLENCE & THREATS

Workers reviewed passages involving violence and threatening language.

HATE SPEECH

Workers labeled text containing hateful or abusive material.

SEXUAL ABUSE

Some passages involved sexual abuse and other disturbing content.

NOT JUST CLICKING.

A label requires a decision about meaning and context. A word or phrase can appear in a news report, a personal account, a threat, or an abusive message. Human reviewers were asked to distinguish harmful material so the tool could learn what to flag.

1 **READ**
Review a text passage

2 **INTERPRET**
Consider meaning and context

3 **LABEL**
Assign a category

Simplified illustration of the labeling process; not an actual leaked task or worker interface. (Perrigo, 2023)

WHO PAYS FOR AI SAFETY?

The reported figures reveal a compensation question - but they measure different things.

$12.50

PER HOUR

Contract rate specified for OpenAI payments to Sama.

CLIENT -> SERVICE PROVIDER

$1.32-$2

PER HOUR

Reported worker take-home estimate, varying by role.

WORKER PAY (TIME ESTIMATE)

DO NOT READ THESE AS A PROFIT CALCULATION.

The $12.50 was a vendor contract rate; the other figure was estimated worker take-home pay. The contract rate also covered operating and employment-related costs. The difference does not tell us Sama's profit.

THE ACCOUNTS ALSO DIFFER.

WORKERS TOLD TIME

Three interviewed workers described reviewing roughly 150-250 passages during a nine-hour shift.

SAMA RESPONDED

Sama disputed the workload account, saying workers were assigned about 70 passages per shift.

Sama also disputed TIME's pay estimates, stating workers could earn $1.46-$3.74 per hour after taxes. These are competing accounts, not a single settled wage figure. (Perrigo, 2023)

WORKERS REPORTED PSYCHOLOGICAL HARM.

WORKERS' ACCOUNTS

All four workers interviewed by TIME described being mentally scarred by the work. They said counseling was rare or unhelpful; two said they were offered group sessions only.

SAMA + OPENAI RESPONDED

Sama said workers could access individual and group sessions with licensed therapists. OpenAI said it understood that counseling, opt-out without penalty, and exposure limits were available; it said Sama managed pay and mental-health provisions.

These are competing accounts, not a complete independent assessment of every worker's experience. All figures describe the historical 2021-2022 project. (Perrigo, 2023)

WHERE DOES THE VALUE GO?

The project separates the people doing the work from the organizations that commission and manage it.

ANNOTATORS

DO THE WORK

Read, interpret, and label harmful text. Receive wages for their labor.

What they contribute: human judgment.

SAMA

EMPLOYS + MANAGES

Contracts with the client, organizes the workforce, and receives the service payment.

What it controls: employment arrangements.

OPENAI

COMMISSIONS + USES

Pays for the outsourced work and uses the resulting safety capability in its AI product.

What it controls: the product and its integration.

THE INEQUALITY IS ABOUT MORE THAN WAGES.

Outsourcing can separate the place where work is performed from the place where product decisions are made. It can also make responsibility harder to trace: the developer commissions a service, the vendor manages employment, and the worker has less control over the final product.

THE MARKET PRESSURE BEHIND OUTSOURCING

Muldoon and colleagues' fieldwork at three Sama delivery centres in Kenya and Uganda argues that competition in the outsourcing market can pressure providers to limit social commitments in favor of higher profit margins. Their analysis describes a model in which value from low-paid data work can be extracted for investors. This is a finding about Sama's wider business model - not a calculation of profit from the OpenAI contract. (Muldoon et al., 2025)

WHAT THIS CASE DOES - AND DOES NOT - SHOW

The OpenAI project shows different actors controlled different parts of the work. Public evidence does not establish Sama's exact profit on this contract or quantify every party's financial gain. It does show why AI developers should treat labor standards as part of their own supply-chain responsibility. (Perrigo, 2023; Muldoon et al., 2025)

HACK THE SYSTEM THAT HIDES ITS WORK.

Mar Hicks shows how technical systems can encode social assumptions and distribute recognition, rights, and power.

HICKS'S HISTORICAL CASE

Hicks examines mid-twentieth-century British government systems that relied on fixed gender classifications. Those assumptions affected transgender citizens' records, benefits, employment, and pay. Their efforts to challenge the bureaucracy made the system's hidden rules visible.

SYSTEM CATEGORIES

AI safety labels are not neutral facts. People define categories and make judgments about what counts as harmful, safe, or acceptable.

WHO BECOMES VISIBLE?

The label enters the data and becomes legible to the model. The person who made the judgment and the conditions of that work can disappear from view.

THE LABEL BECOMES VISIBLE TO THE MACHINE. THE WORKER DISAPPEARS FROM THE PRODUCT STORY.

This is a comparison of how systems encode assumptions and distribute recognition - not a claim that the historical experiences Hicks studies are the same as the experiences of outsourced data workers.

WHAT DOES “HACKING” LOOK LIKE HERE?

Worker testimony, investigative reporting, independent labor assessments, and enforceable sourcing standards challenge the story that AI safety is purely automated. Making the labor visible is a first step toward changing the system. (Hicks, 2019; Perrigo, 2023; Partnership on AI, 2022)

WHAT WOULD ACCOUNTABLE AI LABOR REQUIRE?

Ethical AI should include the people who make its systems usable - not only the users who see the finished product.

1

01 FAIR PAY

Pay workers fairly, with transparent rates and a living-wage benchmark.

2

02 SAFE CONDITIONS

Assess exposure to harmful content; provide appropriate support and a genuine ability to opt out.

3

03 FAIR CONTRACTS

Use clear terms, predictable pay, and protections against unpaid time.

4

04 RESPONSIBLE MANAGEMENT

Set realistic workloads, communicate changes, and provide independent channels to raise concerns.

5

05 WORKER VOICE

Support worker representation and include workers in decisions about the conditions of their labor.

Standards draw on Fairwork's five principles and Partnership on AI's data-enrichment guidelines (Cant et al., 2023; Partnership on AI, 2022).

IF AI IS BUILT BY HUMANS,
ITS PROGRESS MUST NOT HIDE THEIR COST.

READ THE EVIDENCE.

The claims in this zine are grounded in the sources below.TIME documents the specific OpenAI-Sama project; the academic and institutional sources provide wider context and standards.

1. Perrigo, B. (2023, January 18).

Exclusive: OpenAI used Kenyan workers on less than $2 per hour to make ChatGPT less toxic. TIME. https://time.com/6247678/openai-chatgpt-kenya-workers/

2. Muldoon, J., Cant, C., Graham, M., & Ustek Spilda, F. (2025).

The poverty of ethical AI: Impact sourcing and AI supply chains. AI & Society, 40, 529-543. https://doi.org/10.1007/s00146-023-01824-9

3. Cant, C., Ustek Spilda, F., Brittain, L., & Graham, M. (2023).

Fairwork AI Ratings 2023: The Workers Behind AI at Sama. Global Partnership on AI. https://fairwork.oii.ox.ac.uk/wp-content/uploads/sites/11/2023/12/Fairwork-AI-Ratings-2023_final.pdf

4. Partnership on AI. (2022).

Data Enrichment Sourcing Guidelines. https://partnershiponai.org/wp-content/uploads/2022/11/data-enrichment-guidelines.pdf

5. International Labour Organization. (2024, March 20).

Digital labour platforms in Kenya: Exploring women's opportunities and challenges across various sectors. https://researchrepository.ilo.org/esploro/outputs/report/Digital-labour-platforms-in-Kenya-exploring/995366289002676

6. Hicks, M. (2019).

Hacking the cis-tem: Transgender citizens and the early digital state. IEEE Annals of the History of Computing, 41(1), 20-33. https://doi.org/10.1109/MAHC.2019.2897667

IMAGE, AI & PRODUCTION CREDITS

Cover collage: AI-generated conceptual illustration created with ChatGPT image generation; it does not depict actual workers or a documented workplace. Diagrams and page layouts: original project graphics. Research synthesis and editing: Student; factual claims checked against the cited sources. PDF layout: Python / ReportLab. Student: Le Anh Tu.